

# Rule Mining in Knowledge Graphs

January 24, 2025

# Motivation for Knowledge Graphs

*The vision of Semantic Web : enable algorithms to interpret, reason, and interact intelligently with information.*

How to represent knowledge in a form that can be processed by machines?

# How to represent knowledge?

## **Relational Model:**

- Well-defined structure, efficient for tabular data.
- Typically used for restrained number of relations.
- However, real-life relations are in the order of tens of thousands.

## **RDF: Resource Description Framework**

- A standard proposed by the W3C to structure knowledge.
- Flexible and extensible format, suited for open and interconnected systems.

# Principle of the RDF Model

## Describe knowledge as elementary assertions:

*Subject — Predicate — Object*

### Example:

- Marie Curie **was born in** Warsaw.
- Marie Curie **was married to** Pierre Curie.
- Marie Curie **lived in** Paris.
- Marie Curie **is a** scientist.
- A scientist **is a** person.

Simple model, but several considerations must be addressed: ambiguities, consistency, and predicate interoperability.

# Addressing ambiguity

## Representing knowledge with unique identifiers (IRIs) :

- For **predicates** (relations).
- For **entities** (people, concepts, places).

# Addressing ambiguity

## Representing knowledge with unique identifiers (IRIs) :

- For **predicates** (relations).
- For **entities** (people, concepts, places).

## In Turtle, our examples can be expressed as follows :

@prefix yago: <http://yago-knowledge.org/resource/>.

@prefix schema: <https://schema.org/> .

@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>.

@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#>.

yago:Marie_Curie	schema:birthPlace	yago:Warsaw.
yago:Marie_Curie	schema:marriedTo	yago:Pierre_Curie.
yago:Marie_Curie	schema:livedIn	yago:Paris.
yago:Marie_Curie	rdf:type	yago:Scientist.
yago:Scientist	rdfs:subclassOf	yago:Person.

## Textual descriptions: multilingual and multi-type

- Motivation: Some values need to be represented directly.
- A **literal** is a textual or numerical value considered to be fixed.
- A literal may be associated with a **language tag** or a **data type**.

### Example :

```
yago:Marie_Curie rdfs:label "Marie Curie"@fr.  
yago:Marie_Curie rdfs:label "Maria Skłodowska-Curie"@pl.  
yago:Marie_Curie schema:birthDate "1867-11-07"^^xsd:date.
```

**Note :** The predicate `rdfs:label` is generally used to associate an entity with its **textual label** :

```
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#>
```

# Knowledge Graphs Overview

Intuitively, a knowledge graph is a set of *triples*:

*Subject — Predicate — Object*


**Definition:** Let the following sets be given:

- $\mathcal{E}$ : the set of entities (IRI).
- $\mathcal{P}$ : the set of predicates (IRI).
- $\mathcal{L}$ : the set of literal values (e.g., strings, numbers).

A knowledge graph is a set  $\mathcal{K} \subseteq (\mathcal{E} \cup \mathcal{P}) \times \mathcal{P} \times (\mathcal{L} \cup \mathcal{E})$ .



# Simplified Graph Representation



## Simplifications for readability:

- We omit the prefixes of IRI identifiers.
- We use images for some entities instead of their IRIs.
- We do not explicitly represent the labels (`rdf:label`).

# Rules and Integrity Constraints

## Examples of natural language expressions:

- Married people live in the same city.
- A person has only one birth date.
- Every subject of the predicate `marriedTo` is of type `Person`.
- Two authors of the same publication with the same name are the same person.

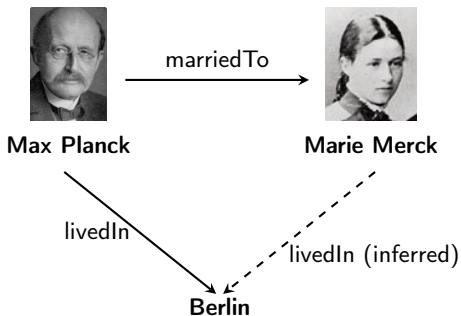
# Rules and Integrity Constraints

## Examples of natural language expressions:

- Married people live in the same city. *(Horn rule)*
- A person has only one birth date. *(functional dependency)*
- Every subject of the predicate `marriedTo` is of type `Person`.  
*(inclusion dependency)*
- Two authors of the same publication with the same name are the same person.  
*(key constraint)*

# Applications of Rules: Fact Prediction

**Rule:** Married people live in the same locality.




# Type Inference

**Rule:** Every subject of the `marriedTo` relation is of type `Person`.

马克斯·普朗克  $\xrightarrow{\text{marriedTo}}$  玛丽·默克

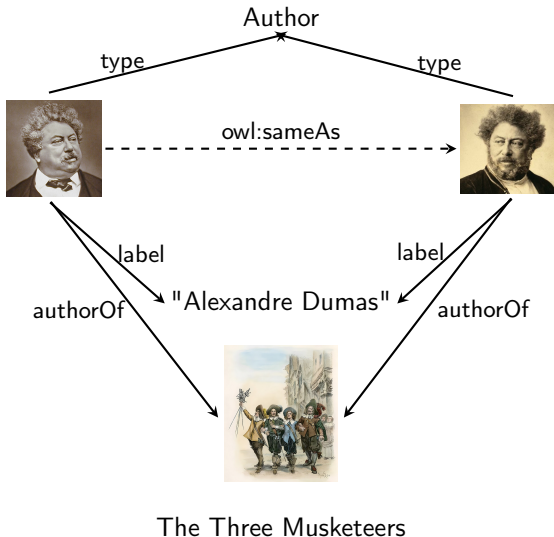
# Type Inference

**Rule:** Every subject of the `marriedTo` relation is of type `Person`.



## Inferring equivalent entities

**Rule:** Two authors sharing the same name and the same publications are the same person.



## Reinforcing constraints

**Rule:** A person can have only one birthdate.



**Exception raised:** two birthdates for the same person.

**Applications:** detecting inconsistencies and cleaning data.



# How to learn rules from Knowledge Graphs?

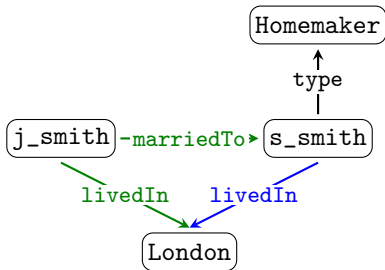

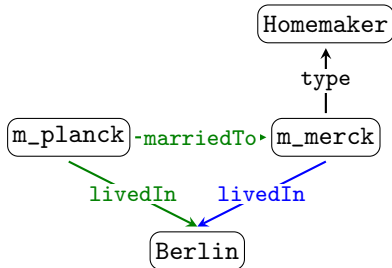

**Induction approach:** Generalise a number of similar observations into a hypothesis.

**Example:** Given many examples of spouses living together, generalise this knowledge.



## Intuition

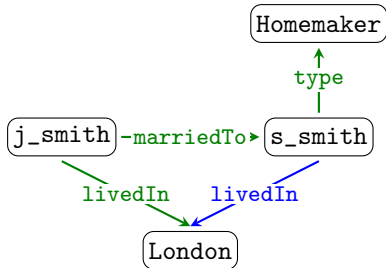
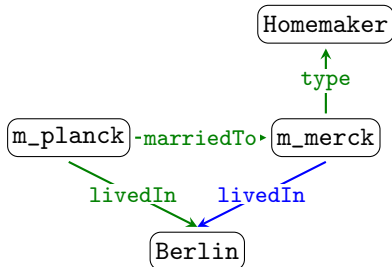


### Knowledge Base :



**Observation:** In 75% of the cases, spouses are recorded as living together.

## Reducing the scope of the analysis

### Knowledge Base:



**Observation:** Reducing the analysis to spouses where one is a homemaker, the rule that the two live together holds in 100% of cases.

# Challenges

- How to find the right balance between generality and specificity?

# Challenges

- How to find the right balance between generality and specificity?
- How to evaluate the rules?

# Challenges

- How to find the right balance between generality and specificity?
- How to evaluate the rules?
- Knowledge bases are incomplete.  
How can we distinguish between false and missing information?

# Challenges

- How to find the right balance between generality and specificity?
- How to evaluate the rules?
- Knowledge bases are incomplete.  
How can we distinguish between false and missing information?
- The search space is too vast to cover efficiently.

# Challenges



- How to find the right balance between generality and specificity?
- How to evaluate the rules?
- Knowledge bases are incomplete.  
How can we distinguish between false and missing information?
- The search space is too vast to cover efficiently.
- The search space depends on the expressivity of the rule language.



# Challenges

- How to find the right balance between generality and specificity?
- How to evaluate the rules?
- Knowledge bases are incomplete.  
How can we distinguish between false and missing information?
- The search space is too vast to cover efficiently.
- The search space depends on the expressivity of the rule language.
- We need smart heuristics to uncover good rules.

## Conjunctive queries for extracting frequent patterns




**Query:** Find all pairs (`?y`, `?z`) such that the spouse of `?y` lives in `?z`.



The query in SPARQL :

```
SELECT ?y ?z
WHERE {
  ?x marriedTo ?y.
  ?x livedIn ?z.
}
```

The query as a graph pattern:



## Conjunctive queries for extracting frequent patterns



### Answers


?x	?y	?z
m_curie	p_curie	Paris
m_planck	m_merck	Berlin

**Query:** Find all pairs (?y, ?z) such that the spouse of ?y lives in ?z.

The query in SPARQL :



```
SELECT ?y ?z
WHERE {
  ?x marriedTo ?y.
  ?x livedIn ?z.
}
```

The query as a graph pattern:



## Datalog rewriting of the same query

### Knowledge Graph:



## Datalog rewriting of the same query

### **Knowledge Graph as a set of facts (grounded atoms):**

<code>marriedTo(m_curie, p_curie)</code>	<code>marriedTo(m_planck, m_merck)</code>
<code>livedIn(m_curie, Paris)</code>	<code>livedIn(m_planck, Berlin)</code>

## Datalog rewriting of the same query

### Knowledge Graph as a set of facts (grounded atoms):

marriedTo(m_curie, p_curie)	marriedTo(m_planck, m_merck)
livedIn(m_curie, Paris)	livedIn(m_planck, Berlin)

### Datalog Query:

$$\text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$$

## Datalog rewriting of the same query

### Knowledge Graph as a set of facts (grounded atoms):

marriedTo(m_curie, p_curie)	marriedTo(m_planck, m_merck)
livedIn(m_curie, Paris)	livedIn(m_planck, Berlin)

### Datalog Query:

$$\text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$$

- $\text{marriedTo}(x, y), \text{livedIn}(x, z)$ , and  $R_1(y, z)$  are **atoms**.

## Datalog rewriting of the same query

### Knowledge Graph as a set of facts (grounded atoms):

<code>marriedTo(m_curie, p_curie)</code>	<code>marriedTo(m_planck, m_merck)</code>
<code>livedIn(m_curie, Paris)</code>	<code>livedIn(m_planck, Berlin)</code>

### Datalog Query:

$$\text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$$

- `marriedTo(x, y)`, `livedIn(x, z)`, and  $R_1(y, z)$  are **atoms**.
- The variables  $x$ ,  $y$ , and  $z$  are placeholders for values from the KG.



## Datalog rewriting of the same query

### Knowledge Graph as a set of facts (grounded atoms):

<code>marriedTo(m_curie, p_curie)</code>	<code>marriedTo(m_planck, m_merck)</code>
<code>livedIn(m_curie, Paris)</code>	<code>livedIn(m_planck, Berlin)</code>

### Datalog Query:

$$\text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$$

- `marriedTo(x, y)`, `livedIn(x, z)`, and  $R_1(y, z)$  are **atoms**.
- The variables  $x$ ,  $y$ , and  $z$  are placeholders for values from the KG.
- $R_1$  is the **query result relation**, generated from the rule.

## Datalog rewriting of the same query

### Knowledge Graph as a set of facts (grounded atoms):

<code>marriedTo(m_curie, p_curie)</code>	<code>marriedTo(m_planck, m_merck)</code>
<code>livedIn(m_curie, Paris)</code>	<code>livedIn(m_planck, Berlin)</code>

### Datalog Query:

$$\text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$$

- `marriedTo(x, y)`, `livedIn(x, z)`, and  $R_1(y, z)$  are **atoms**.
- The variables  $x$ ,  $y$ , and  $z$  are placeholders for values from the KG.
- $R_1$  is the **query result relation**, generated from the rule.

**Query Answers:** Each valid **substitution** for  $x$ ,  $y$ , and  $z$  that satisfies the rule generates a row in the result:

?x	?y	?z
m_curie	p_curie	Paris
m_planck	m_merck	Berlin

# Introduction to Horn Rules

**Example Rule  $R$ :**  $\text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livedIn}(y,z)$

- A fact  $p$  predicted by the Knowledge Graph  $\mathcal{K}$  and rule  $R$  is denoted:

$$\mathcal{K} \wedge R \models p$$

Here,  $\models$  indicates logical entailment, meaning  $p$  is implied by  $\mathcal{K}$  and  $R$ .

# Introduction to Horn Rules

**Example Rule  $R$ :**  $\text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livedIn}(y,z)$

- A fact  $p$  predicted by the Knowledge Graph  $\mathcal{K}$  and rule  $R$  is denoted:

$$\mathcal{K} \wedge R \models p$$

Here,  $\models$  indicates logical entailment, meaning  $p$  is implied by  $\mathcal{K}$  and  $R$ .

- The rule states: If a person  $x$  is married to  $y$  and  $x$  lived in a city  $z$ , then  $y$  is also inferred to have lived in  $z$ .

# Introduction to Horn Rules

**Example Rule  $R$ :**  $\text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livedIn}(y,z)$

- A fact  $p$  predicted by the Knowledge Graph  $\mathcal{K}$  and rule  $R$  is denoted:

$$\mathcal{K} \wedge R \models p$$

Here,  $\models$  indicates logical entailment, meaning  $p$  is implied by  $\mathcal{K}$  and  $R$ .

- The rule states: If a person  $x$  is married to  $y$  and  $x$  lived in a city  $z$ , then  $y$  is also inferred to have lived in  $z$ .
- The variables  $x$ ,  $y$ , and  $z$  are **universally quantified**, meaning the rule applies to all individuals in  $\mathcal{K}$ .

# Introduction to Horn Rules

**Example Rule  $R$ :**  $\text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livedIn}(y,z)$

- A fact  $p$  predicted by the Knowledge Graph  $\mathcal{K}$  and rule  $R$  is denoted:

$$\mathcal{K} \wedge R \models p$$

Here,  $\models$  indicates logical entailment, meaning  $p$  is implied by  $\mathcal{K}$  and  $R$ .

- The rule states: If a person  $x$  is married to  $y$  and  $x$  lived in a city  $z$ , then  $y$  is also inferred to have lived in  $z$ .
- The variables  $x$ ,  $y$ , and  $z$  are **universally quantified**, meaning the rule applies to all individuals in  $\mathcal{K}$ .

## Application to our KG:

- $\mathcal{K} \wedge R \models \text{livedIn}(\text{p\_curie}, \text{Paris})$
- $\mathcal{K} \wedge R \models \text{livedIn}(\text{m\_merck}, \text{Berlin})$

## Mining Horn Rules: Reducing the Search Scope

- Systems like **AMIE** and **AnyBURL** focus on mining **positive Horn rules**:
  - The conclusion consists of a single positive literal.
  - The premise is a conjunction of positive literals.
- To further reduce the search space, they impose additional constraints:
  - **Connectedness**: All atoms in the rule must be transitively connected.
  - **Closedness**: A rule is closed if:
    - All variables appear in at least two atoms.
- A closed rule ensures **safely**: variables in conclusion appear in the premise.

# Expressive Power

**Question:** Among the following formulas, which are not Horn rules?

- Married people live in the same city.
- A person has only one date of birth.
- Every subject of the predicate `marriedTo` is of type `Person`.
- Two authors of the same publication with the same name are the same person.



# Expressive Power

**Question:** Among the following formulas, which are not Horn rules?

- **Married people live in the same city.**

$\text{marriedTo}(x, y), \text{livesIn}(x, z) \Rightarrow \text{livesIn}(y, z)$



- **A person has only one date of birth.**
- **Every subject of the predicate `marriedTo` is of type `Person`.**
- **Two authors of the same publication with the same name are the same person.**

# Expressive Power

**Question:** Among the following formulas, which are not Horn rules?

- **Married people live in the same city.**

$\text{marriedTo}(x, y), \text{livesIn}(x, z) \Rightarrow \text{livesIn}(y, z)$



- **A person has only one date of birth.**

$\text{birthdate}(x, d_1), \text{birthdate}(x, d_2) \Rightarrow d_1 = d_2$



- **Every subject of the predicate marriedTo is of type Person.**
- **Two authors of the same publication with the same name are the same person.**

# Expressive Power

**Question:** Among the following formulas, which are not Horn rules?

- **Married people live in the same city.**

$\text{marriedTo}(x, y), \text{livesIn}(x, z) \Rightarrow \text{livesIn}(y, z)$



- **A person has only one date of birth.**

$\text{birthdate}(x, d_1), \text{birthdate}(x, d_2) \Rightarrow d_1 = d_2$



- **Every subject of the predicate marriedTo is of type Person.**

$\text{marriedTo}(x, y) \Rightarrow \text{type}(x, \text{'Person'})$



- **Two authors of the same publication with the same name are the same person.**

## Expressive Power

**Question:** Among the following formulas, which are not Horn rules?

- **Married people live in the same city.**

$\text{marriedTo}(x, y), \text{livesIn}(x, z) \Rightarrow \text{livesIn}(y, z)$  ✓

- **A person has only one date of birth.**

$\text{birthdate}(x, d_1), \text{birthdate}(x, d_2) \Rightarrow d_1 = d_2$  ✗

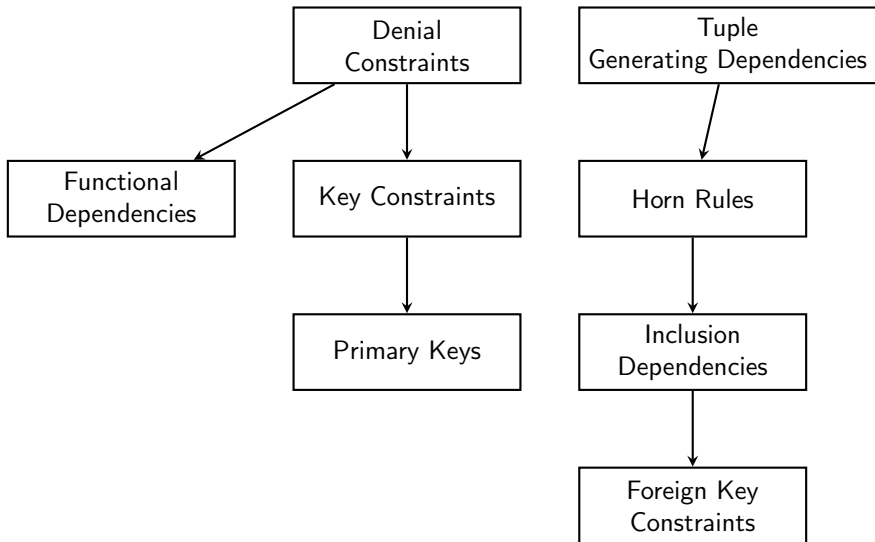
- **Every subject of the predicate `marriedTo` is of type `Person`.**

$\text{marriedTo}(x, y) \Rightarrow \text{type}(x, \text{'Person'})$  ✓

- **Two authors of the same publication with the same name are the same person.**

$\text{authorOf}(x, p), \text{authorOf}(y, p), \text{label}(x, n), \text{label}(y, n) \Rightarrow x = y$  ✗

## Expressivity: Hierarchy of Languages



# Lattice: Rule Refinements Exploration

**Initial Query:**  $\Rightarrow \text{livesIn}(y,z)$

$\Rightarrow \text{livesIn}(y,z)$

*Add Dangling Atom*




$\text{marriedTo}(x,y) \Rightarrow \text{livesIn}(y,z)$

*Add Closing Atom*



$R_1: \text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livesIn}(y,z)$

*Add Instantiated Atom*



$R_2: \text{marriedTo}(x,y), \text{livedIn}(x,z), \text{type}(x, \text{'Homemaker'}) \Rightarrow \text{livesIn}(y,z)$

# Query Containment

Rules can be generated by **extending existing rules**, forming a containment relationship between rules.

## Example:

$R_1 : \text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livedIn}(y,z)$

$R_2 : \text{marriedTo}(x,y), \text{livedIn}(x,z), \text{type}(x, \text{'Homemaker'}) \Rightarrow \text{livedIn}(y,z)$

- Does  $R_2$  generate a subset of predictions compared to  $R_1$ ?

# Inclusion of Conjunctive Queries

## Queries:

$R_1 : \text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$

$R_2 : \text{marriedTo}(x, y), \text{livedIn}(x, z), \text{type}(x, \text{'Homemaker'}) \Rightarrow R_2(y, z)$

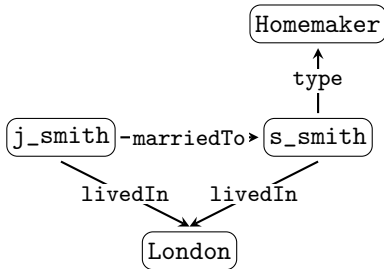



- Regardless of the knowledge base, the results of query  $R_2$  are always included in those of query  $R_1$ .
- Formally,  $R_2$  is a sub-query of  $R_1$ , denoted as:

$$R_2 \sqsubseteq R_1$$



# KG Example

## Knowledge Base:



# Specialization of Rules

The specialization of a rule can be achieved in two main ways:

- **Adding atoms** to the body of a rule (as seen previously).
- **Replacing variables with constants**, restricting the application of the rule to a specific subset.

## Example:

$$R_1 : \text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow \text{livedIn}(y, z)$$
$$R_3 : \text{marriedTo}(x, y), \text{livedIn}(x, \text{'Paris'}) \Rightarrow \text{livedIn}(y, \text{'Paris'})$$

**Observation:**  $R_3$  applies only to people married to Parisians.

# Specialization of Rules

The specialization of a rule can be achieved in two main ways:

- **Adding atoms** to the body of a rule (as seen previously).
- **Replacing variables with constants**, restricting the application of the rule to a specific subset.

**Example:** Let us compare the results of the two rules:

$$R_1 : \text{marriedTo}(x, y), \text{livedIn}(x, z) \Rightarrow R_1(y, z)$$


$$R_3 : \text{marriedTo}(x, y), \text{livedIn}(x, \text{'Paris'}) \Rightarrow R_3(y, \text{'Paris'})$$

Notice that here as well,  $R_3 \sqsubseteq R_1$ .

**Observation:**  $R_3$  applies only to people married to Parisians.

# Rule Generation Lattice

- Rule generation can be represented as a lattice, where rules are connected by the inclusion operator  $\sqsubseteq$ .
- General rules are at the top, while specific rules are at the bottom.
- Lattices optimize rule generation by factoring computations and avoiding redundancies.



**Question:** Where should rule specialization stop?

**Challenge:** Finding the right balance between generality and specificity.

# Rule Support

**Definition:** The support of a rule  $R$  is the number of true facts it generates:

$$\text{support}(R) = |\{p : (\mathcal{K} \wedge R \models p) \wedge p \in \mathcal{K}\}|$$

**Property:** The support decreases as a rule becomes more specialized.

**Utility:** A rule is not refined if its support falls below a threshold.

**Question:** What support threshold should you choose for your data?

# Calculating Rule Support

## Rules:

$R_1 : \text{marriedTo}(x,y), \text{livedIn}(x,z) \Rightarrow \text{livedIn}(y,z)$

$R_2 : \text{marriedTo}(x,y), \text{livedIn}(x,z), \text{type}(x, \text{'Homemaker'}) \Rightarrow \text{livedIn}(y,z)$

## Predictions:

$R_1$

livedIn	
p_curie	Paris
m_merck	Berlin
f_faure	Paris
s_smith	London

$R_2$

livedIn	
m_merck	Berlin
s_smith	London

## Confidence Measure of a Rule

**Definition:** The confidence of a rule  $R$  is the proportion of true predictions among all predictions:

$$\text{confidence}(R) = \frac{\text{support}(R)}{\text{support}(R) + |\text{cex}(R)|}$$

where  $\text{cex}(R)$  represents the counterexamples of  $R$ .

**Extension:** AMIE 3 generalizes Inductive Logic Programming (ILP) by discovering *soft* rules that tolerate a limited number of counterexamples.

### Summary:

- **Support:** Relevance of a rule.
- **Confidence:** Accuracy of a rule.

**Objective:** Extract rules with support and confidence above defined thresholds.

# The Problem of Counterexamples

**Observation:** Knowledge bases primarily contain positive information but lack explicit counterexamples.

**Solution:** Negative facts can be inferred based on assumptions:

- **Closed World Assumption (CWA):** What is not known is false.
- **Open World Assumption (OWA):** What is not known is unknown (and could therefore be true).
- **Partial Completeness Assumption (PCA):** Proposed by AMIE.

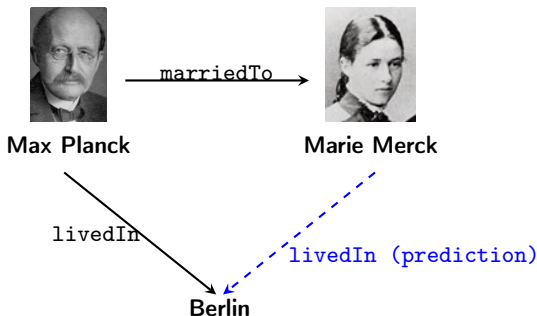


# Counterexample or True Prediction?

Consider the prediction:

$$\mathcal{K} \wedge R \models \text{livedIn}(\text{Marie Merck}, \text{Berlin})$$

where  $\mathcal{K}$  is the following knowledge base:



## Interpretation based on assumptions:

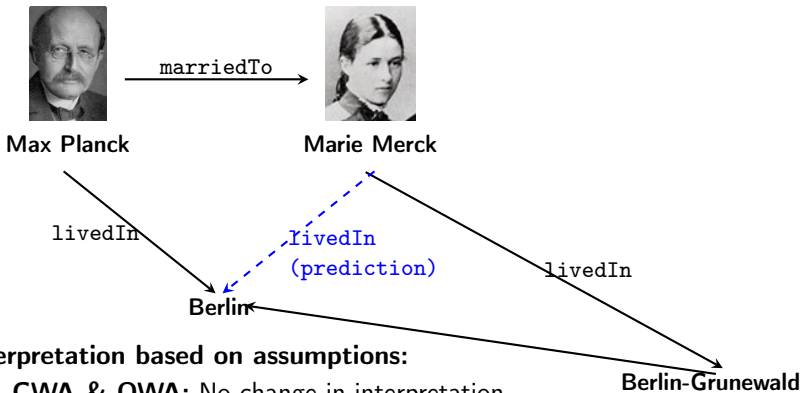
- **CWA:** The prediction is false, so it is a counterexample.
- **OWA:** No conclusion can be drawn; the fact might be unknown.
- **PCA:** The prediction is true because no other residence is known for Marie Merck.

# Knowledge Base Update

Consider the prediction:

$$\mathcal{K}' \wedge R \models \text{livedIn}(\text{Marie Merck}, \text{Berlin})$$

where  $\mathcal{K}'$  is the following knowledge base:



**Interpretation based on assumptions:**

- **CWA & OWA:** No change in interpretation.
- **PCA:** The prediction is a counterexample, as another fact  $\text{livedIn}(\text{Marie Merck}, \text{Berlin-Grunewald})$  is known.

# Overview of Rule Discovery Methods

- **Inductive Logic Programming (ILP):** Finding hypotheses that cover examples.
- **Approaches:**
  - *Top-Down:* Based on specialization (e.g., AMIE, RUDIK).
  - *Bottom-Up:* Based on generalization (e.g., GOLEM, AnyBURL).
- **Interesting Perspective:**
  - Rule discovery on vector representations.

# Bibliography

- Luis Galárraga, Christina Teflioudi, Katja Hose, Fabian M. Suchanek: *“AMIE: Association Rule Mining under Incomplete Evidence”* WWW 2013.
- N. Lavrač, S. Džeroski: *“Inductive Logic Programming — Techniques and Applications.”* References.
- F. Suchanek, J. Lajus, A. Boschini, G. Weikum: *“Knowledge Representation and Rule Mining in Entity-Centric KBs”* Reasoning Web Summer School (RW), 2019.